

Implementing Baker’s SUBTYPEP decision procedure

Léo Valais
Jim E. Newton
Didier Verna
lvalais@lrde.epita.fr
jnewton@lrde.epita.fr
didier@lrde.epita.fr
EPITA/LRDE
Le Kremlin-Bicêtre, France

ABSTRACT

We present here our partial implementation of Baker’s decision procedure for subtypep. In his article “A Decision Procedure for Common Lisp’s SUBTYPEP Predicate”, he claims to provide implementation guidelines to obtain a subtypep more accurate and as efficient as the average implementation. However, he did not provide any serious implementation and his description is sometimes obscure. In this paper we present our implementation of part of his procedure, only supporting primitive types, CLOS classes, member, range and logical type specifiers. We explain in our words our understanding of his procedure, with much more detail and examples than in Baker’s article. We therefore clarify many parts of his description and fill in some of its gaps or omissions. We also argue in favor and against some of his choices and present our alternative solutions. We further provide some proofs that might be missing in his article and some early efficiency results. We have not released any code yet but we plan to open source it as soon as it is presentable.

CCS CONCEPTS

• **Theory of computation** → **Type theory**; *Divide and conquer*; Pattern matching.

ACM Reference Format:

Léo Valais, Jim E. Newton, and Didier Verna. 2019. Implementing Baker’s SUBTYPEP decision procedure. In *Proceedings of the 12th European Lisp Symposium (ELS’19)*. ACM, New York, NY, USA, 8 pages.

1 INTRODUCTION

The Common Lisp standard [1] provides the predicate function subtypep for introspecting the sub-typing relationship. Every invocation (subtypep A B) either returns the values (t t) when A is a subtype of B, (nil t) when not, or (nil nil) meaning the predicate could not (or failed to) answer the question. The latter can happen when the type specifier (satisfies P) (representing the type $\{x \mid P(x)\}$ for some predicate and total function¹ P) is involved.

¹A function defined over its entire definition domain.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

ELS’19, April 01–02 2019, Genova, Italy

© 2019 Copyright held by the owner/author(s).

ACM ISBN 978-2-9557474-3-8.

```
(sb!xc:deftype keyword ()  
  '(and symbol (satisfies keywordp)))
```

Listing 1: The keyword type definition in SBCL

For example, given two arbitrary predicates F and G, there is no way subtypep can answer the question (subtypep '(satisfies F) '(satisfies G)).

However, some implementations abuse the permission to return (nil nil). For example, in SBCL 1.4.10 (the implementation we are currently focusing our efforts on), (subtypep 'boolean 'keyword) returns (nil nil), thus violating the standard². The definition of the keyword type is responsible for this failure: as shown in Listing 1, it involves a satisfies type specifier³.

Another kind of problem for which subtypep’s accuracy matters is the optimization of the typecase construct as shown in [7] and [8]. The aim is to remove redundant checks in the construct and the approach is to use binary decision diagrams. However, to build such a structure, subtypep is repeatedly used. The unreliability of the predicate leads here to many lost BDD reductions and therefore to the generation of sub-optimal code.

Our implementation is still in active development, currently targets SBCL and focuses almost entirely on result accuracy. It supports *primitive* types, *user-defined* types (deftype, classes and structures), member (and eql) type specifiers and ranges (e.g., (integer * 12)). We present our strategy for implementing each one of these while discussing how and why we decided or not to diverge from Baker’s [3] approach—or potentially filling some gaps or unclear bits. No optimization work has been done yet and the implementation still has bugs and diverse issues, but we have found some encouraging results about accuracy and even about efficiency.

2 THE COMMON LISP TYPE SYSTEM

2.1 Type specifiers

Common Lisp types are not manipulated directly. Instead, the type to be manipulated is *described* using a *type specifier*. The type specifier Domain-Specific Language (DSL) allows programmers to describe types by writing S-expressions which obey some rules described in the Common Lisp standard [1].

²The Common Lisp standard requires that no invocation of subtypep involving only primitive types return (nil nil).

³C.f. bug #1533685 in SBCL bug tracker.

```
(deftype except (x)
  `(not (eq1 ,x)))
```

Listing 2: The deftype construct

A subtlety about type specifiers is that *different ones* can represent the *same* type (e.g., `integer`, `(integer * *)` and `(or fixnum bignum)` all describe the same type). This means that *symbolic computation* does *not suffice* to answer the sub-typing question. Note that one could write a predicate, say `type=`, to determine whether two type specifiers in fact describe the same type using two calls of `subtypep`.

It is possible to define *parametric aliases* using the `deftype` construct. It is then possible to refer to a whole type specifier using its alias. Listing 2 shows an example of parametric `deftype`.

2.2 Vocabulary

type	A set of elements. For any type $u: u \equiv \{x \mid x : u\}$
canonical t.s.	A type specifier without aliases.
primitive type	A standardized type ([2]) that is not necessarily implemented as a class.
symbolic form	A type specifier whose type is <code>symbol</code> .
compound form	A type specifier whose type is <code>list</code> .
logical form	A compound form whose car is <code>or</code> , <code>and</code> or <code>not</code> .
kingdom	In Baker's terminology, a "type kingdom" designates the types that can be described using <i>only one kind of type specifier</i> . <code>nil</code> (the empty type) belongs to every type kingdom.

In this article we focus on two particular type kingdoms:

- the *literal type kingdom*, represented using only symbolic, member and logical type specifiers, and,
- the *range type kingdom*, represented only using range and logical type specifiers

For example, `(or string symbol)` belongs to the literal type kingdom. `(and number (not real))` belongs to the range type kingdom. However, `(or symbol integer)` belongs to the literal type kingdom while `(or symbol (integer * *))` belongs to both. This situation is handled in section 4.

There are other type kingdoms that Baker mentions in his article, such as the *array* type kingdom, represented using only `array` and logical type specifiers. Note that a type can belong to several kingdoms, as multiple type specifiers can describe it. For example, `integer` belongs to literal *and* range kingdoms as the type specifiers `integer` (symbolic) and `(integer * *)` (range) both describe it. In Section 4, we describe how to guarantee that a given type is only described by one kind of type specifier, hence restricting it to *one* kingdom.

3 PROCEDURE'S MECHANISMS OVERVIEW

Figure 1 shows the internals of our implementation. Every step will be detailed in the following sections. There are three *major stages*:

- (1) *The pre-processing* — Both type specifiers are processed in order to simplify further calculations: the aliases are expanded, and each occurrence of numeric types are converted to their

equivalent range type specifier. Finally, as explained there-after, the procedure splits into several sub-procedures, one for each type kingdom, because their internal type representation differ. In order to achieve that, the type specifiers must also be split into equivalent subtype specifiers restricted to each concerned kingdom. This stage is detailed in Section 4.

- (2) *Expert sub-procedures* — Once split, each subtype specifier is redirected to the appropriate expert sub-procedure. The job of such a procedure is to *prove*, in its own kingdom, the assertion "*A* is a subtype of *B*" to be *wrong*. Our procedures currently only support literal and range type specifiers—an expert sub-procedure has been implemented only for these two kingdoms. This stage is detailed in Section 5.
- (3) *Result conjunction* — Eventually, all expert sub-procedures return (a Boolean) and the results are accumulated using conjunction. (In practice, as soon as one expert procedure returns false, `subtypep` returns.)

4 PRE-PROCESSING

4.1 Alias expansion

The very first step is to ensure that the type specifier is in its *canonical form*, that is, having all its aliases expanded. This is done by the `expand` function. For example, considering the type created in Listing 2, `(expand '(except 12))` should return `(not (eq1 12))`.

Unlike macro expansion, `deftype` expansion is not standardized in Common Lisp. Thus a solution must be found for each Common Lisp implementation independently. As our efforts are currently focused on SBCL, we discuss how we implement the `expand` function for that compiler.

SBCL's `subtypep` heavily relies on the function `sb-kernel:specifier-type`, which does type expansion. It also does type simplification—turning `(and integer string)` into `nil`—which could have saved us some work. We hoped we could simplify that function to make it compatible with Baker's algorithm while keeping the `deftype` expansion and the range canonicalization work. However we found, thanks to [7] tools, that the function is responsible for most of the work of `subtypep`, as shown in Figure 2. Considering the lack of efficiency of that function and the fact that it would not be trivial to simplify it to only keep the interesting bits, we decided on another, more cost-effective solution.

The function `sb-ext:typexpand` takes a type specifier and tries to expand it (not recursively). It either returns the expansion result, or the input type specifier if it is not expandable. `(sb-ext:typexpand 'integer)` returns `integer` since it is not a `deftype` alias whereas `(sb-ext:typexpand '(except 12))` returns `(not (eq1 12))`. To expand a whole type specifier, it just needs to walk through it, applying `sb-ext:typexpand` on each list or atom manually. One subtlety though is that the result of an expansion may itself be an alias to expand⁴. For example, let's say that we have `(deftype my-type () '(except 0.0))`, then the result of `(sb-ext:typexpand 'my-type)` is `(except 0.0)`, which is of course an alias to expand again.

⁴Fortunately, `sb-ext:typexpand` also returns a Boolean indicating whether or not an expansion happened.

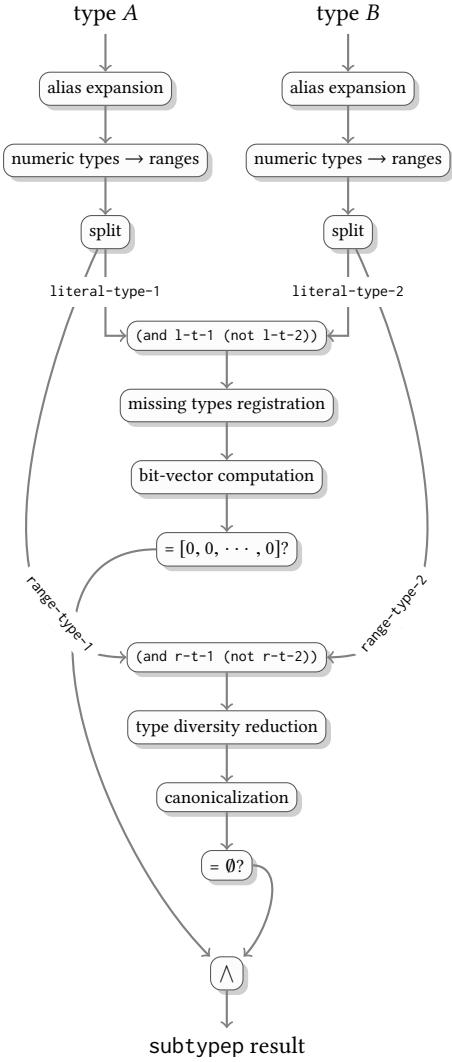



Figure 1: Internal flowchart of (subtypep A B)

Figure 2: specifier-type weight in c1:subtypep execution^a

^a cached-subtypep-caching-call is just a memoizing wrapper around SBCL’s subtypep which is a bit more efficient than the raw implementation.

4.2 Numeric type specifiers conversion

As explained in Section 3, after pre-processing both type specifiers, the procedure splits in two expert sub-procedures: one for literal type specifiers and one for range type specifiers. Numeric types—types containing numbers (mathematically speaking)—can have different representations: a *symbol* (e.g., `fixnum`), a *member expression* (e.g., `(member 1 2 3)`) or a *range* (e.g., `(integer 1 6)`). However, the first two belong to the literal type kingdom whereas the latter belongs to the range kingdom. Thus, the numerical type information would be distributed over the different expert sub-procedures. For consistency and accuracy, a single internal representation has to be chosen. The symbolic and member numeric types must each be converted into an equivalent type specifier, in which numerical data are only represented using ranges.

- *Symbolic* numeric type specifier — say `U`, replace it by `(U * *)`⁵. Note the new “type specifier” is likely *not* to be valid (e.g., `(fixnum * *)` is invalid). Because it is never exposed to the user—as it is an intermediate, internal representation—nothing bad can happen. However, it cannot be used with other functions requiring a type specifier, such as `typep`.
- *member* type specifiers — e.g., `(member a 1 2 :b)` is converted to `(or (member a :b) (bit 1 1) (integer 2 2))`. To do that,
 - (1) extract the numbers out of the expression,
 - (2) map each number, say `n`, to construct the type specifier `((type-of n) n n)`⁶,
 - (3) and combine the remaining member expression and the ranges with the `or` logical type specifier.

A subtlety to consider is that *super-types* of number also contain numerical data that must be extracted. Indeed, the type atom contains both numerical data—`(number * *)`—and non-numerical data—`(and atom (not (number * *)))`. Thus, its replacement in the numeric type kingdom is straightforward: `(number * *)`. In the literal type kingdom however, its replacement is `(or stream array character function standard-object symbol structure-object structure-class)`. The type `t`—which is `(or atom sequence)`—must be converted similarly.

Yet another subtlety is that the type specifiers `(and)` and `(or)` respectively describe the types `t` and `nil`. Hence every occurrence of `(and)` must be replaced by the replacement of `t` described in the previous paragraph. In order to remove that annoying corner case completely, `(or)` is also replaced, by `nil`.

4.3 Splitting

Having reached this step, the input now only contains *canonical literal* and *range* type specifiers, numeric types being *only* expressed as ranges. The next stage—expert sub-procedures—requires literal and numeric types to be separated.

Thus the top type `t` is divided into two⁷ disjoint subtypes—“kingdoms” as Baker says. The previous step, described in Section 4.2, ensures that the representation (in terms of type specifiers) of the types in each kingdom is different. All numeric types are

⁵Implementations supporting the IEEE floating point raise many concerns with `-0.0`, `NaN`, `+∞` and `-∞`. Baker explains in detail how to handle these cases.

⁶The results of `type-of` are implementation-dependent. We suppose here that `type-of` only returns the name (as a symbol) of the type of `n` (`n` being a number).

⁷One per kingdom actually, but since our implementation only supports two—literal and range types—we only focus our attention on these.

represented as ranges, and literal types as symbolic and member (without numbers) type specifiers.

This step roughly consists of an in-depth traversal of the type specifier, using pattern-matching to recognize which type specifier represents which type. We use the implementation of [9] because of its simplicity and versatility.

Our implementation uses a function `type-keep-if` which takes a predicate P and a type specifier U and returns:

- U as it is when $P(U) = \top$,
- nil when $P(U) = \perp$,
- $(\text{op } U_1 \dots U_n)$ where $U_i = (\text{type-keep-if } P \ U_i)$ when $U = (\text{op } U_1 \dots U_n)$ and $\text{op} \in \{\text{and, or, not}\}$.

Given the predicate `literal-type-p` and a type U , `type-keep-if` returns U with every inner type specifier that describes a non-literal type replaced by `nil` (interpreted as the *empty type*). The result is then a subtype of $(\text{and } (\text{not number}) (\text{not } (\text{array } * *)))$. Likewise, given the predicate `range-type-p`, this function returns U with every non-range inner type specifier replaced by `nil` (interpreted this time as the *empty range*). Thus, the result is a subtype of number. Therefore, `split` can easily be implemented in terms of `type-keep-if`.

4.4 Type reformulation

For any types U and V , $U \subseteq V \Leftrightarrow U \cap \bar{V} = \emptyset$. Therefore, for any type specifiers U and V , when `(subtypep U V)` returns \top , then `(subtypep '(and ,U (not ,V)) nil)` also returns \top .

The results of the `split` function are zipped together using `(lambda (x y) '(and ,x (not ,y)))` before being passed to the expert sub-procedures. This way, they will not have to prove that an arbitrary type is a subtype of *another* arbitrary subtype, but rather whether *one* arbitrary type specifier describes the empty type (which is substantially easier to reason about, and implement).

5 EXPERT SUB-PROCEDURES

Listing 3 shows how `subtypep` could be defined from a top-down point of view. It shows that, according to Figure 1, both type specifiers are processed independently, split into two kingdoms (literal and numeric types) and unified in an `(and U (not V))` fashion. The expert sub-procedures, `null-literal-type-p` and `null-numeric-type-p`, each accept one argument—a type specifier, say U —and returns a Boolean indicating whether U describes the empty type (`nil`).

Each sub-procedure answers restricted to its kingdom—as no type can (at this point of the procedure) belong to two different kingdoms, as shown in section 4. With that piece of information, we can (now) safely assert that:

- the literal type kingdom is the type described by `(and (not number) (not (array * *)))`⁸, and,
- the numeric type kingdom is the type described by `number`⁹.

⁸Actually this is not completely accurate since the type string can be described using array type specifiers. However, since the latter are not supported by our implementation yet, we consider the types `string` and `bit-vector` as being literal types since their symbolic representation is kept through the entire process. This is very likely to change in the future.

⁹Our implementation does not support complex numbers yet, and considers the complex type as being empty. Some wrong results arise from that supposition—such as `(subtypep`

```
(defun subtypep (a b)
  (reduce (lambda (x y) (and x y))
    (mapcar (lambda (expert t1 t2)
      (funcall expert `(and ,t1 (not ,t2))))
    (list #'null-literal-type-p
      #'null-numeric-type-p)
    (split (num-types->ranges (expand a))
      (split (num-types->ranges (expand b))))))
```

Listing 3: A top-down approach of `subtypep`

There are several properties that are derived from the preceding pre-processing steps. First of all, both kingdoms' procedures are guaranteed to only ever receive argument *canonical* type specifiers. These are also guaranteed to never contain `atom` or `t` type specifiers. The occurrences of `(and)` and `(or)` have been replaced respectively by `t` and `nil`. `eql` type specifiers have been replaced by equivalent member expressions. `member` type specifiers only occur in the literal type kingdom and contain no numerical data. Numerical data are only expressed as intervals, which are likely not to be valid type specifiers. Both kingdoms accept the type specifier `nil` but with a *different meaning*: for literal types, `nil` means the empty type which complement is `t` whereas for numeric types it represent the empty range whose complement is `(number * *)`.

In the following sections we describe in detail the implementation of the expert sub-procedures for the literal (Section 5.1) and numeric (Section 5.2) type kingdoms. We also briefly discuss in Section 5.3 the array type kingdom and the `cons` type specifier family, which Baker ignores in his article.

5.1 Procedure for literal types

5.1.1 Theory. To represent types in the literal types kingdom, we suppose at first that there is a way to enumerate every element in t , say $e_1, e_2, \dots, e_\omega$. Then, let $u_1, u_2, \dots, u_\omega$ be all the (non-strict) subtypes of the top-level type t . We associate to each pair (u_i, e_j) the bit b_{ij} with the value 1 when $e_j \in u_i$ and 0 when $e_j \notin u_i$. Let bv_i be the *representative bit-vector* associated to the type u_i , defined by $[b_{i0}, b_{i1}, \dots, b_{i\omega}]$. These bit-vectors are the rows of the infinite matrix on Eq. $B_{\omega\omega}$ which illustrates the system.

$$\begin{array}{c}
 \begin{matrix} & e_1 & e_2 & e_3 & e_4 & \cdots & e_\omega \\
 u_1 & \left(\begin{array}{cccccc} 1 & 0 & 0 & 0 & \cdots & 1 \end{array} \right) \\
 u_2 & \left(\begin{array}{cccccc} 0 & 1 & 1 & 0 & \cdots & 0 \end{array} \right) \\
 u_3 & \left(\begin{array}{cccccc} 0 & 0 & 0 & 1 & \cdots & 0 \end{array} \right) \\
 \vdots & \left(\begin{array}{cccccc} \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \end{array} \right) \\
 u_\omega & \left(\begin{array}{cccccc} 1 & 0 & 1 & 0 & \cdots & 0 \end{array} \right)
 \end{matrix}
 \end{array}
 \quad (B_{\omega\omega})$$

PROOF. Each type has a unique bit-vector representation.

Let u_i and u_j be two distinct types. Thus, $(u_i \cup u_j) \setminus (u_i \cap u_j) \neq \emptyset$. Let $e_k \in (u_i \cup u_j) \setminus (u_i \cap u_j)$. By definition, we have $b_{ik} \neq b_{jk}$. Hence $bv_i \neq bv_j$. Two distinct types are represented by two different bit-vectors.

Similarly, let bv_i and bv_j be two different bit-vectors. Then it necessarily exists a k such as $b_{ik} \neq b_{jk}$. Therefore $\exists e_k, (e_k \notin u_i \vee e_k \notin u_j) \wedge e_k \notin u_i \cap u_j$. Hence $u_i \neq u_j$. \square

¹⁰`'number 'real)` returning true. This will change as soon as complex numbers are supported.

PROOF. Type intersection, union and complement are equivalent to bitwise Boolean operations “and”, “or” and “not” on representative bit-vectors.

Let two types u_i and u_j in:

- (1) Let $u_k = u_i \cup u_j$. By definition, $\forall l \in \mathbb{N} \cup \{\omega\}, b_{kl} = 1$ iff $b_{il} = 1$ or $b_{jl} = 1$, that is $b_{kl} = b_{il} \vee b_{jl}$. Thus, also by definition:

$$\begin{aligned} bv_k &= [b_{k0}, b_{k1}, \dots, b_{k\omega}] \\ &= [b_{i0} \vee b_{j0}, b_{i1} \vee b_{j1}, \dots, b_{i\omega} \vee b_{j\omega}] \\ &= bv_i \vee bv_j \end{aligned}$$

- (2) We proceed similarly for the intersection and the Boolean logical operator “and” (\wedge).
- (3) Let $u_k = \overline{u_i}$. We have by definition $\forall l \in \mathbb{N} \cup \{\omega\}, b_{kl} = \neg b_{il}$. Then:

$$\begin{aligned} bv_k &= [\neg b_{i0}, \neg b_{i1}, \dots, \neg b_{i\omega}] \\ &= \neg bv_i \end{aligned} \quad \square$$

5.1.2 Implementation. Common Lisp cannot enumerate all the possible subtypes of t nor all of its elements. Fortunately, we do not need them all. We only need to consider the types mentioned in the input type specifier to determine its emptiness.

We also do *not need to enumerate all the elements of these types*. It is that aspect of the procedure of Baker that makes it both powerful and difficult to understand at first. We only need *sufficiently many* elements from a type to distinguish it from the other types. Because we are now considering only a *finite* number of types, say u_1, \dots, u_n , to register a new type u_{n+1} to our (now *finite*) matrix, we only need to find an element $e \in u_{n+1}$ such as $e \notin u_1 \cup \dots \cup u_n$.

Now let's suppose that the type specifier of u_{n+1} is in fact (member e), that e is itself chosen as a representative element for another type, say u_k , and that u_k is only distinguished from the other registered types by that element e . u_{n+1} and u_k would then have the same bit-vector representation when these types are likely to be distinct. The general solution for that kind of problem is to register *all* the elements found inside the member type specifier. When there is a conflicting element e already registered as a representative for another types, we generate additional representatives for these types. That precaution ensures that this kind of conflict never happens and greatly simplifies the implementation of member type specifiers.

To implement that registration matrix system, we use two functions: B : type name \rightarrow bit-vector, with $B(u_i) = bv_i$, and I : representative \rightarrow bit index, with $I(e_i) = i - 1$. Baker suggests in his small example [3] using the operator set which is deprecated in modern Common Lisp programming. Instead, we use hash tables to represent these functions. Type names are symbols, bit-vectors are bit-vectors and element indexes are positive integers. To register a new type u_{n+1} , it is added to the B hash table and its bit-vector content $b_{(n+1)i}$ is evaluated for all the existing representatives ($i \in [1; m]$). To register a new representative e_{m+1} , it is added to the I hash table with the index m . Then we add one bit (the m -th bit) to each bit-vector bv_i and evaluate it in respect to the type u_i . Thus, to retrieve the bit-vector of a registered primitive or user-defined type t , we just lookup its value $B(t)$. To compute the bit-vector of a member expression (member $e_1 \dots e_n$), we use the

value $B(\text{member } e_1 \dots e_n) = \bigvee_{i=1}^n \beta(I(e_i))$, where $\beta(x)$ returns the null bit-vector with the x -th bit activated.

The bit-vector of logical type specifiers are given in Eq. 1, Eq. 2 and Eq. 3 thereafter.

$$B((\text{and } U_1 \dots U_n)) = \bigwedge_{i=1}^n B(U_i) \quad (1)$$

$$B((\text{or } U_1 \dots U_n)) = \bigvee_{i=1}^n B(U_i) \quad (2)$$

$$B((\text{not } U)) = \neg B(U) \quad (3)$$

5.1.3 Issues. The method for choosing the representative elements for a type depends of its nature: it can be a primitive type, a user-defined type (class, structure or condition) or a member expression.

Since primitive types are known (c.f. table 4.2 of [2]), their representative elements are chosen at compile-time. The u_{n+1} subtlety above should still be kept in mind. For instance, the type `null` is a subtype of both `symbol` and `list`; so three representative elements are needed: `nil`, a non-empty list and a symbol other than `nil`. Note that some primitive types are an *exhaustive partition* of other types (e.g., `character` \equiv (or `base-char` `extended-char`)). Obviously, in that case, such a precaution does not apply.

For user-defined types, Baker suggests to extend the type creation mechanism—thus modifying the implementation's internal functions—to register a dummy element as a representative. We decided *not* to follow his approach because of the poor portability of his solution. Indeed, this work, often non-trivial, would have to be repeated for each targeted Common Lisp implementation. (We would like to avoid modifying the SBCL internal mechanisms.) Moreover, it would register a representative for *every* class created, thus increasing bit-vectors' size uselessly since only a few of these classes are likely to appear in a subtypep type specifier. But more importantly, the main drawback of his solution is that creating that dummy element might have unexpected side-effects, as it may need to use slot's default values and/or `initialize-instance`. We decided instead to use the *Meta Object Protocol* (MOP) [6], more specifically *class prototypes*. Class prototypes are pseudo-instances of a class, created without executing `initialize-instance` and which typep and `eq` view as traditional instances. However, to create a class prototype, the class needs to be *finalized* and it cannot be guaranteed until it is instantiated. Since that class may be involved in a subtypep call before that happens, when a new class is encountered, we force its finalization using the function `ensure-finalized` from the (portable) `closer-mop` package¹⁰. Then, we create the prototype of the class using `sb-mop:class-prototype` and register it. This method is much more portable than Baker's and does not require to hook inside the implementation.

Since (in SBCL¹¹) conditions are classes, they are supported automatically. The Common Lisp standard [1] states that “`defstruct` without a `:type` option defines a class with the structure name as its name”, hence in that case no additional work is required. The standard also states that “Specifying this option [...] prevents the

¹⁰<http://common-lisp.net/project/closer/>

¹¹Every major lisp implementations implement conditions as CLOS classes—the most obvious way to do it. We ignore exotic condition implementations.

structure name from becoming a valid type specifier recognizable by `typep`.” Thus, `subtypep` is not concerned by these types of structures.

To address the misrepresentation problem when `member` type specifiers are involved, as discussed in Section 5.1.3, we must ensure that a new representative element is generated and registered. The Common Lisp standard ([1]) states that the `member` type specifier is defined in terms of `eql`. That is, `(typep e '(member e1 ... en))` uses `eql` to compare `e` to the successive `ek` to check the membership. That precise property reduces the misrepresentation problem to only two types: `symbol` and `character` (and their subtypes).

To better understand why it is the case, first consider a reduced version of the top-level type `t`: `t = (or string list symbol)`. Then, let `R = ("hello" (1 2 3) foo)` be our list of representatives.

- (1) Let's ask the question `(subtypep 'symbol '(member foo))`.
- (2) As discussed in Section 5.1.3, we add the elements of the `member` expression to `R`. To conform with the specification, we first check whether or not `foo` is already in `R` `eql`-wise: `foo ∈eql R`, so `R` does not change.
- (3) As shown in Eq. 4, the emptiness check passes, meaning that `symbol` is indeed a subtype of `(member foo)`, which is obviously wrong.

$$\begin{aligned} B(\text{symbol}) \wedge \neg B(\text{member } \text{foo}) &= 001 \wedge \neg 001 & (4) \\ &= 001 \wedge 110 \\ &= 000 \\ &= B(\text{nil}) \end{aligned}$$

However, for lists, that problem does not appear, thanks to the `eql`-wise comparison.

- (1) `(subtypep 'list '(member (1 2 3)))`
- (2) `(1 2 3) ∉eql R ⇒ R = ("hello" (1 2 3) foo (1 2 3))`
- (3) As shown in Eq. 5, the emptiness check fails and the answer is correct.

$$\begin{aligned} B(\text{list}) \wedge \neg B(\text{member } (1\ 2\ 3)) &= 0101 \wedge \neg 0001 & (5) \\ &= 0101 \wedge 1110 \\ &= 0100 \\ &\neq B(\text{nil}) \end{aligned}$$

Within the literal types kingdom, the only types for which this problem occurs—since the representatives are not supposed to be accessible to the user of `subtypep`—are then `symbol` and `character`. Therefore, only the representatives of these types need to be actually checked when registering `member`'s elements.

To generate a new `symbol`, we use `alexandria:symbolicate`¹². The keyword `subtype` of `symbol` is also subject to the problem. (Actually, solving the problem for keywords also solves the problem for symbols.) To generate a new `character`, we first need to know whether it is a `base-char` or an `extended-char`. Then we pick a `character` of that type not registered yet. When all the `characters` of that type are registered there is nothing to do (since the type is fully represented in the matrix, no misinterpretation can occur).

We have not addressed the problem of a type specifier involving a user class `C` and a `member` expression containing the class prototype of `C` yet.

5.2 Procedure for numeric types

Unlike the literal type kingdom, the range type kingdom does not need an internal state to represent numeric types. Indeed, the expert sub-procedure takes as input an already precise enough representation of the type described. Range type specifiers allow to describe which kind of number is specified (its type, e.g., integer, ratio, etc.), its bounds (inclusive and exclusive, e.g., `(integer (0) 6)`) and is able to represent non-bounded intervals through the symbol `*` meaning infinity (e.g., `(float * 0.0) ≡ [-∞; 0.0]`). The range type specifier is as precise as the mathematical range notation. Additionally, the mathematical union, the intersection and complement of these ranges can be expressed equally using the corresponding logical type specifier.

Therefore, to assert about the emptiness of the input type specifier, checking whether the *canonicalized* version of this interval expression describes the empty range (i.e., `nil`) is *sufficient*. The calculation is performed by three successive steps, which we describe in the following sections.

This algorithm suffers from an exponential time and space complexity. However, Baker claims that in practice, that theoretical complexity is not an issue (it only appears for “highly artificial cases”). We have not tried to prove (or invalidate) his statement but Section 6 shows some early results that tend to support his claim.

We use a custom abstraction, the interval class, closer to the mathematical object (with type, bounds and limits slots). Thus we avoid the annoying manipulation of lists (with the many standardized ranges syntaxes). The first step is to write a function `range->interval` that converts (using pattern matching) a range type specifier to its corresponding interval instance. This function also takes care of the exotic compound forms—such as `(unsigned-byte s)` which describes the integer range `[0; 2s - 1]`. We also use a similar structure for interval operations to fully discard the list representation.

We also need the following interval functions:

- `(interval-and I1 I2)` — returns `I1 ∩ I2` if their types are `eql`, or `∅` otherwise.
- `(interval-or I1 I2)` — returns `I1 ∪ I2` if their types are `eql` and `I1 ∩ I2 ≠ ∅`, or `∅` otherwise.
- `(interval-minus I1 I2)` — returns `I1 - I2` (may return two values when `I2 ⊂ I1`) if their types are `eql`, or `I1` otherwise.
- `(interval-empty-p I)` — returns whether `I = ∅`.

5.2.1 Type diversity reduction. Functions working with intervals must be aware of the relationship of the types of these intervals. For example, the intersection of two integer intervals might be non-empty whereas the intersection of one integer and one single-float intervals is always null as these two types are disjoint. However, integer and `fixnum` are different types but the intersection of intervals of such types *might be non-empty*. The subtype relationship of the types of intervals needs to be introspected to accurately apply some operations (such as intersection or union).

The type number (complex numbers being ignored) is an exhaustive partition of six mutually disjoint types: integer, ratio, single-float, short-float, double-float, and long-float. Baker advises to define what he calls “simple intervals”, that is intervals guaranteed to have their type equal to one of these six types. This

¹²<https://common-lisp.net/project/alexandria/>

<i>Supertype</i>	<i>Conversion</i>
number	(or rational float)
real	(or rational float)
rational	(or integer ratio)
float	(or short-float single-float double-float long-float)
bignum	(or integer (not fixnum))

Table 1: Conversion table for supertypes

way, as these types are mutually disjoint, operations on intervals of such types have their implementation greatly simplified.

To convert each numeric type into its equivalent using only the six types above, a two-step conversion is required.

- (1) For intervals whose type is a *supertype* of one of these types, the conversion table 1 is used. E.g.: the conversion of (rational $a\ b$) gives (or (integer $a\ b$) (ratio $a\ b$)).
- (2) For intervals whose type is a *bounded subtype* (i.e.: having defined bounds, not infinity) of these six types, their actual bounds have to be constrained to fit within the bounds of their type, before being converted to their corresponding supertype. For example, (fixnum $12\ 2^{100}$), has to be converted to (integer most-negative-fixnum most-positive-fixnum), where most-positive-fixnum $< 2^{100}$, as 2^{100} is a bignum, thus discarding the numbers in between. A similar procedure is applied to the types bit, short-float, single-float, double-float and long-float.

Eventually, the type of every interval is constrained to one of the six types above, with the bounds (if some) of their original type preserved.

5.2.2 Canonicalization. To check the emptiness of the interval expression, it is canonicalized. Let Γ be the canonicalization function. Its parameter is either an interval I or an operation on intervals χ (intersection, union or complement). Γ either returns \emptyset , an interval or a *union of disjoint intervals*—the three possible outcomes of a mathematical interval canonicalization.

First and foremost, anytime Γ encounters or returns a union, it must ensure that it is *flattened* (no nested unions). It must also ensure that the intervals inside the union are *disjoint*. As shown in Section 5.2.1, intervals with different types are necessarily disjoint. *Touching intervals* [3] are *merged* using interval-or.

$\Gamma(\emptyset)$ and $\Gamma(I)$ are straightforward, as shown in Eq. end- \emptyset and Eq. end- I . These are the terminal cases of the recursion of Γ .

$$\Gamma(\emptyset) = \emptyset \quad (\text{end-}\emptyset)$$

$$\Gamma(I) = I \quad (\text{end-}I)$$

Intersections (and logical type specifiers) are reduced as soon as they are encountered. Their operands need to be processed by Γ first (hence the implicit mapping “ $k \rightarrow n$ ”). Eq. and-apply shows how to reduce intersections. The Φ_f operator denotes a *fold* [5] operation using the function f . $\Gamma \circ \cap$ denotes the composition of the Γ function and the intersection operator. To break it down in a bottom-up fashion:

- (1) Eq. and-final — the application of the intersection function.

- (2) Eq. and-distribution’ — the distribution of the intersection over the union. Next step is Eq. and-final.
- (3) Eq. and-distribution — also the distribution of the intersection over the union. However, $\Gamma(\chi)$ may return an union, leading the execution either to Eq. and-distribution’ or directly to Eq. and-final.
- (4) Eq. and-apply — the canonicalization of the χ_n forms using mapping. The results are then folded using $\Gamma \circ \cap$, thus initiating the recursive intersection distribution.

$$\Gamma\left(\bigcap_n \chi_n\right) = \Phi_{\Gamma \circ \cap} \Gamma(\chi_k)_{k \rightarrow n} \quad (\text{and-apply})$$

$$\Gamma\left(\chi \cap \bigcup_n I_n\right) = \bigcup_n \Gamma(\chi \cap I_n) \quad (\text{and-distribution})$$

$$\Gamma\left(\bigcup_n I_n \cap I\right) = \bigcup_n \Gamma(I_n \cap I) \quad (\text{and-distribution'})$$

$$\Gamma(I_1 \cap I_2) = (\text{interval-and } I_1\ I_2) \quad (\text{and-final})$$

Complements (not logical type specifiers) are also reduced as soon as they are encountered. Their only operand is first canonicalized. Complementing U in number (the top-level type of the range type kingdom) is equivalent to the difference number $-U$, as shown in Eq. not-apply. The difference canonicalization goes through a similar recursive distribution path than the intersection, that is Eq. minus-distribution and then Eq. minus-apply. Note that this path is taken every time since the interval difference is an internal operation and that its left-hand operand is always \mathcal{U} .

$$\Gamma(\bar{\chi}) = \Gamma(\mathcal{U} - \Gamma(\chi)) \quad (\text{not-apply})$$

$$\mathcal{U} = \langle \text{type diversity reduction of (number * *)} \rangle$$

$$\Gamma\left(\chi - \bigcup_n I_n\right) = \bigcup_n \Gamma(\chi - I_n) \quad (\text{minus-distribution})$$

$$\Gamma\left(\bigcup_n I_n - I\right) = \bigcup_n (\text{interval-minus } I_n\ I) \quad (\text{minus-apply})$$

5.2.3 Range emptiness check. Once an interval expression canonicalized, checking its emptiness is trivial. The predicate interval-empty-p, given the result of the first Γ call, just returns the Boolean that null-numeric-type-p has to return.

5.3 Array types and cons type specifiers

This section presents some preliminary work and research results found on array and cons type specifiers. Obviously, since the implementation of the expert sub-procedures for these kingdoms is still a work in progress, no result nor implementation guidelines are provided here. It does, however, give some insights about how Baker procedure applies to modern Common Lisp implementations such as SBCL.

Array type specifiers are complex to handle because they are bi-dimensional: it has an element type and bounds (e.g., (array integer (* 2 *))). Internally, Common Lisp implementations do not store which exact type specifier is specified but rather only store

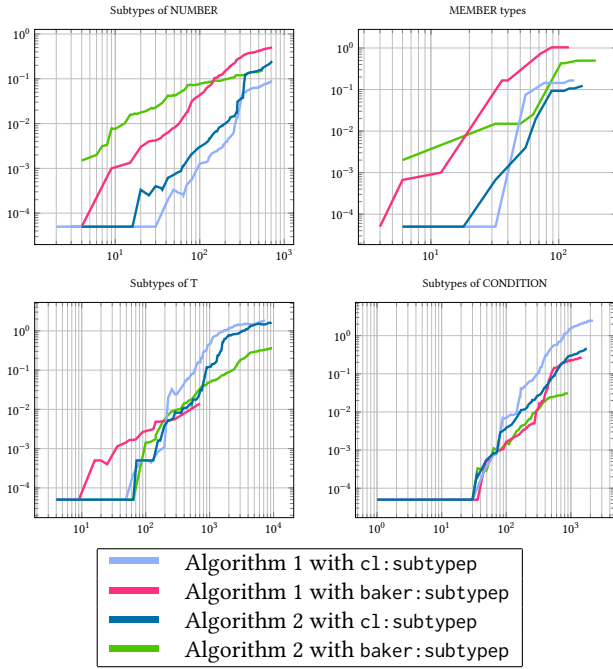


Figure 3: Comparative efficiency measures of our subtypep implementation

the result of the function `upgraded-array-element-type` returns giving that type. E.g, for `(make-array 2 :element-type 'list)`, the implementation does not makes an array of `list` but rather an array of `(upgraded-array-element-type 'list)`. For every value that might return this function, Baker requires that we store a bit matrix (instead of bit vectors) because of the complex bounds logic of the type specifier. As for the literal type procedure, it seems to be an efficient type representation system—albeit more complex—which nonetheless requires an extra registration step and a global state.

Baker does not mention the `cons` type specifier family at all in his article because it appeared after he released his article [4]. An accurate expert sub-procedure for this kingdom would have an exponential complexity. More investigation is needed to assert whether or not that exponential time is “acceptable” (as it is for ranges) before rejecting it. The accuracy of existing subtypep procedures for the `cons` type specifier also needs to be studied.

6 EARLY RESULTS

Our implementation of subtypep is still in active development and very experimental. No serious optimization work has been made. Nonetheless, Newton has compared in [7] the performances of several subtypep highly dependent algorithms, both using the implementation of SBCL and ours.

These results, shown in Figure 3, are only presented here as complementary information. On the horizontal axis is the size of the type specifiers and on the vertical axis is the measured execution time. Hence, the lower a curve is, the better. As expected, our implementation is often slower, but not dramatically, which is encouraging.

- Our implementation is overall slower in the range type kingdom.
- Heavy users of `member` seems to experience a slower execution. Perhaps, as predicted by Baker, the reason is that the systematic registration of the elements makes the size of the bit-vectors grow quickly, thus making every subsequent operation slower.
- For the symbolic type specifiers—primitive types, `CLOS` classes and conditions—our implementation already outperforms SBCL's.

7 CONCLUSION AND FUTURE WORK

Throughout this article we presented our implementation of Baker's decision procedure. In Section 2 we introduced the Common Lisp type system, the notion of type specifier and some vocabulary. In Section 4 we explained how to pre-process the caller's type specifiers to make the work of the expert sub-procedures presented in Section 5 easier. We described our implementation for the symbolic, `member`, `range` and logical type specifiers. We also gave some insights about the implementation for the `array` and `cons` type specifiers. We finally presented some early efficiency measures, which are globally encouraging.

Our implementation is still a work in progress and highly experimental. But with some cleaning and the implementation of both `array` and `cons` expert sub-procedures, it could be a viable alternative to existing subtypep implementations. We will have open sourced its code by then. We still have to find a solution for the satisfies type specifier and the related uncertainty. Indeed, in some situations, subtypep still can answer even though the type specifier is involved. For example, in `(subtypep 'string '(and number (satisfies evenp)))`, as the second operand is guaranteed to be a subtype of number, the predicate can safely return false. Finally, a lot of measures on accuracy and efficiency are needed to assert whether Baker's intuition about his procedure was correct or not.

Even if, in the future, we are to conclude that our implementation is less efficient than those which already exists, Baker's algorithm would still likely to improve the predicate's accuracy. Lispers would then have the ability to choose whichever subtypep implementation fits their needs the best.

REFERENCES

- [1] Ansi. American National Standard: Programming Language – Common Lisp. ANSI X3.226:1994 (R1999), 1994.
- [2] Ansi. American National Standard: Programming Language – Common Lisp – Type Specifiers (Section 4.2.3). ANSI X3.226:1994 (R1999), 1994. http://www.lispworks.com/documentation/lw50/CLHS/Body/04_bc.htm.
- [3] Henry G. Baker. A Decision Procedure for Common Lisp's SUBTYPEP Predicate. *Lisp and Symbolic Computation*, 1992.
- [4] Paul F. Dietz. “subtypep tests” discussion on gcl-devel, 2005. <https://lists.gnu.org/archive/html/gcl-devel/2005-07/msg00038.html>.
- [5] Graham Hutton. A tutorial on the universality and expressiveness of fold. *Journal of Functional Programming*, 9(4):355–372, July 1999. URL <http://dblp.uni-trier.de/db/journals/jfp/jfp9.html#Hutton99>.
- [6] Gregor J. Kiczales, Jim des Rivières, and Daniel G. Bobrow. *The Art of the Metaobject Protocol*. MIT Press, Cambridge, MA, 1991.
- [7] Jim Newton. *Representing and Computing with Types in Dynamically Typed Languages*. PhD thesis, Sorbonne Université, Paris, France, November 2018.
- [8] Jim Newton and Didier Verna. Approaches in typecase optimization. In *European Lisp Symposium*, Marbella, Spain, April 2018.
- [9] Peter Norvig. *Paradigms of Artificial Intelligence Programming: Case Studies in Common Lisp*. Morgan Kaufmann, 1992.